

AI • CYBER SECURITY • CLOUD ENGINEER

Navid Ghobadpour

Inspired to build practical AI, cloud and cybersecurity solutions that help make our digital communities safer.

PRODUCT PORTFOLIO / 2026

Sydney, NSW, Australia



I learn by building systems that solve real operational problems

I am an Australian citizen based in Sydney, NSW, with a background spanning civil engineering, cybersecurity, cloud and software development.

- Hands-on builder with a strong interest in practical engineering
- Focused on cybersecurity, cloud platforms and artificial intelligence
- Comfortable moving across infrastructure, frontend, backend and data
- Motivated by continuous learning and complete system delivery



Each learning step unlocked a more complete engineering capability



Networking → operating systems → programming → databases → cloud → AI → security engineering

TAFE NSW turned security theory into practical infrastructure skills

DIPLOMA OF INFORMATION TECHNOLOGY

Cyber Security • TAFE NSW • Completed 2026

WINDOWS SERVER

Administration + hardening

LINUX

Operations + command line

ACTIVE DIRECTORY

Identity + Group Policy

NETWORKING

DNS • DHCP • VPN

CLOUD

Compute • storage • identity

PYTHON

Automation + data handling

SECURITY OPERATIONS

Monitoring + investigation

VIRTUALISATION

Enterprise lab environments

INCIDENT RESPONSE

Triage + remediation

Continuous learning supports every layer of the systems I build

NETWORKING

TCP/IP • DNS • VPN

CYBER SECURITY

Detection • response

CLOUD

AWS • Azure

PROGRAMMING

Python • JavaScript

SOFTWARE ENGINEERING

APIs • architecture

ARTIFICIAL INTELLIGENCE

LLMs • agents

DEVOPS

Git • Docker • CI/CD

OPERATING SYSTEMS

Linux • Windows

DATABASES

SQL • PostgreSQL

Languages become useful when they are applied to complete systems

PYTHON

AI agents • automation • JSON processing

JAVASCRIPT

React frontend • Node.js services

HTML

Semantic application structure

CSS

Responsive layouts • visual systems

SQL

Incident data • relationships • reporting

BASH

Linux operations • repeatable workflows

The frontend turns security data into analyst decisions

REACT

Component-driven analyst workflows

VITE

Fast development and production builds

HTML + CSS

Accessible, responsive structure

AXIOS

Authenticated API communication

RECHARTS

Operational and executive visualisation


LEAFLET

Geospatial threat context

RESPONSIVE DESIGN


Usable across common screen sizes

Secure backend connects identity, APIs and live operations



API design connects authentication, incident workflows, reporting and real-time notifications.

Relational data keeps incidents, users and evidence connected



PostgreSQL + Neon

- Schema design and relationships
- SQL queries for operational views
- Indexes and query optimisation concepts
- Cloud-hosted persistence

Cloud and DevOps turn local builds into accessible systems

SOURCE

Git • GitHub

BUILD

Vite • Node • Docker

DEPLOY

Vercel • Render

DATA

Neon PostgreSQL

CLOUD

AWS architecture


OPERATE

Linux • logs • health checks

CI/CD concepts: repeatable builds, environment configuration and controlled releases.

CORE CERTIFICATIONS


Credentials validate the foundations behind the engineering work



AWS SOLUTIONS ARCHITECT
AWS • 2025 • Architecture



ISC2 CC
ISC2 • 2024 • Security



ETHICAL HACKER
Cisco • 2024 • Assessment



AMD AI ACADEMY
AMD • 2026 • LLM inference



PYTHON PROFESSIONAL
OpenEDG • 2026 • Python



PYTHON OOP
LinkedIn • 2026 • OOP



DEVSECOPS
LinkedIn • 2025 • AppSec

Additional study broadens the toolkit without replacing project evidence

freeCodeCamp

Responsive Web Design

HTML • CSS • accessibility

freeCodeCamp

Relational Database

SQL • Bash • PostgreSQL

freeCodeCamp

JavaScript Algorithms

JavaScript • problem solving

LinkedIn Learning

Front-End Development

Web application foundations

LinkedIn Learning

Windows Server / AD DS

Identity • DNS • administration

Cisco Academy

Python Essentials 1

Programming fundamentals

Microsoft

Azure fundamentals

Compute • storage • identity

Qualys

Vulnerability Management

Assessment • detection • response

WHY I BUILT THESE PROJECTS

The projects began where tutorials stopped

I wanted to understand how complete systems behave - not only how individual technologies work.




Each project adds a new engineering layer: visibility → intelligence → control.

FEATURED PROJECTS


Three platforms connect the full engineering stack



SOC PLATFORM



AI SECURITY OPERATIONS



AGENTOPS CONTROL TOWER

Frontend → backend → data → cloud → AI → governance

Security signals become an operational decision loop


A case study in AI-assisted triage, investigation and response orchestration.

OVERVIEW

One workspace connects context, priority and action

The platform concept combines security telemetry, incident context and AI-assisted recommendations in a human-supervised workflow.

- Unifies analyst context
- Surfaces priority and rationale
- Keeps human approval at critical boundaries
- Captures an auditable response trail



ARCHITECTURE


The architecture separates evidence, reasoning and execution



Python + PydanticAI + MCP + local LLMs



A pragmatic stack supports real-time, secure workflows



FEATURES

Features are designed around analyst trust



EXPLAINABLE PRIORITY

Recommendations include visible context, not black-box output.

HUMAN APPROVAL

High-impact actions remain explicitly controlled.

LIVE CONTEXT

Real-time notifications reduce stale investigation state.

EXECUTIVE TRACE

Operational evidence can flow into clear reporting.

CHALLENGES

The hard problem is safe acceleration—not automation alone

01

OFFLINE INFERENCE

Local Qwen models must deliver useful analysis within constrained environments.

02

MULTI-AGENT COORDINATION

Specialist agents need clear responsibilities and structured handoffs.

03

ANALYST TRUST

MITRE mapping, evidence and remediation must remain reviewable.

```
.conda activate ai-engineer)
(ai-engineer) PS C:\Users\navid\AI-Projects\ai-soc-agent> python -m agents.multi_agent_orchestrator

=== AGENT WORKFLOW ===
→ Log Collection Agent
→ Detection Agent
→ MITRE ATT&CK Agent
→ Threat Intelligence Agent
→ Correlation Agent
→ Severity Escalation Agent
→ Case Management Agent
→ Investigation Agent
=====
[✓] Multi-agent investigation completed
(ai-engineer) PS C:\Users\navid\AI-Projects\ai-soc-agent>
```

✓ Problems
✓ Terminal

RESULTS

The outcome is a demonstrable end-to-end operating model

01

MULTI-AGENT WORKFLOW

02

LOCAL LLM INFERENCE

03

MITRE MAPPING

04

JSON REPORTING

Evidence demonstrated in the project

- Brute force, SQL injection, XSS and API abuse scenarios
- Severity classification and escalation guidance
- Structured incident reports for analyst review
- Command-line investigation assistant

```
(ai-engineer) PS C:\Users\Navid\AI-Projects\ai-soc-agent> python --agent soc_manager_agent

Immediate      Block source IP
Immediate      Review web application logs
Investigation   Review authentication logs
Investigation   Check database logs
Investigation   Verify no unauthorized access occurred
Remediation     Implement parameterized queries
Remediation     Validate user input
Remediation     Deploy WAF protections
Prevention      Conduct security coding review
Prevention      Schedule web application penetration testing

SOC Manager ✓ COMPLETED

1 SOC Decision:
  • Investigate Within 15 Minutes
2 Executive Summary: The case involves an attempted attack detected by the intrusion detection system, which was initiated through a suspicious SQL injection request made to the public-facing login endpoint of a web application. This incident has been evaluated based on available structured data and does not indicate any successful compromise or confirmed threat action.
3 Confirmed Facts:
  • The user 'guest' from source IP address 183.22.55.9 initiated an attempt to access the login page of a web application.
  • The request contained SQL injection syntax, specifically targeting the username parameter in the '/login/username=admin' URL.
  • No known CVE or CISA KIV vulnerabilities were identified related to this incident.
  • There was no evidence suggesting that this attack utilized a zero-day vulnerability.
  • Source IP 183.22.55.9 has been flagged as malicious by reputation systems with an 80% confidence level.
4 Assumptions / Unknowns:
  • The request alone does not confirm the presence of unauthorized access or successful exploitation of any software vulnerabilities.
  • It is unclear if this attack was part of a larger campaign targeting multiple sites.
  • Further investigation into the source IP's history and relationships to other suspicious entities may be necessary.
5 Business Impact:
  • The incident raised concerns about potential security breaches, although no confirmed compromise has been observed.
  • Immediate actions need to be taken to secure the application against future attacks and to prevent unauthorized access or data exposure.
6 Immediate Next Steps:
```

FUTURE ROADMAP

Roadmap: move from assistance to governed autonomy

NOW

Unified context

Human-approved recommendations

NEXT

Evaluation harness

Quality + safety benchmarks

LATER

Bounded actions

Policy-driven automation

SCALE

Learning loops

Measured operational

Terminal output showing a table of investigation cases:

Case ID	Status	Decision	Attack Type	Source IP	User	Events
CASE-20240812087001	Requested	Investigate within 15 Minutes	sql injection	193.22.15.5	guest	26
CASE-20240812087011	Open	Investigate within 15 Minutes	brute force	45.83.12.18	admin	26
CASE-20240812087012	Open	Investigate within 15 Minutes	os	88.12.44.2	guest	17
CASE-20240812087019	Open	Investigate within 15 Minutes	brute force	12.34.56.24	root	15
CASE-20240812088011	Open	Monitor	sql abuse	18.222.44.19	unknown	10
CASE-20240812088017	Open	Monitor	network anomaly	185.22.44.1	unknown	8
CASE-20240812088024	Open	Investigate within 15 Minutes	brute force	48.98.12.18	administrator	8
CASE-20240812088027	Open	Monitor	Suspicious	28.55.44.9	storage.com	7
CASE-20240812088031	Open	Investigate within 15 Minutes	privilege escalation	51.144.22.8	unknown	6
CASE-20240812088040	Open	Investigate within 15 Minutes	network anomaly	11.75.51.4	sa_admin	6
CASE-20240812088043	Open	Investigate within 15 Minutes	brute force	201.8.111.18	admin	6
CASE-20240812088046	Open	Investigate within 15 Minutes	privilege escalation	106.51.100.22	service_m	6
CASE-20240812088052	Open	Investigate within 15 Minutes	network anomaly	192.8.2.46	unknown	7
CASE-20240812088053	Open	Investigate within 15 Minutes	network anomaly	19.8.1.58	unknown	7
CASE-20240812088054	Open	Investigate within 15 Minutes	brute force	45.83.12.18	root	6
CASE-20240812088056	Open	Investigate within 15 Minutes	privilege escalation	18.8.1.15	webapp	6
CASE-20240812088070	Open	Monitor	Suspicious	183.44.12.1	unknown	4
CASE-20240812088073	Open	Monitor	Suspicious	18.8.1.18	server01	4
CASE-20240812088075	Open	Monitor	Suspicious	183.200.18.1	unknown	4

github.com/golden-horizon/agentic-ai-security-operations-platform

Golden Horizon turns telemetry into analyst workflow


A deployed full-stack SOC experience spanning detection, investigation and reporting.

OVERVIEW

The SOC platform is built for the complete incident journey

Golden Horizon simulates how a security operations team monitors, investigates and responds to cyber threats.

- Incident management dashboard
- Investigation center
- Security event monitoring
- Threat hunting and intelligence
- Executive reporting



A clean service boundary keeps the system understandable



TECHNOLOGY STACK

The stack balances developer speed with security controls

FRONTEND

React • Vite • Axios • Recharts • Leaflet

BACKEND

Node.js • Express • Socket.IO

DATA


PostgreSQL • Neon

SECURITY

JWT • bcrypt • Helmet • CORS • MFA


HOSTING

Vercel • Render



FEATURES

Detection and investigation stay connected



MONITOR

Incidents, events and authentication activity

INVESTIGATE

Timeline context and analyst status workflow

HUNT

Search-oriented threat exploration


REPORT

Technical detail and executive summaries

CHALLENGES

Security controls must work without blocking the analyst

- Protecting access with JWT and two-step verification
- Keeping frontend status synchronized after backend updates
- Delivering live alerts without losing investigation context
- Designing dashboards for both analyst depth and executive clarity



ENGINEERING RESPONSE

Explicit identifiers, current-state refresh, layered controls and audience-specific reporting.

RESULTS


The platform demonstrates a credible deployed security product

Implemented evidence

- Full-stack cloud deployment
- PostgreSQL-backed incident data
- Real-time analyst notifications
- Threat detection for common web attack patterns
- Downloadable executive security reports

LIVE

golden-horizon-soc-frontend.vercel.app



Roadmap: deepen detection fidelity and operational resilience

DETECTION


- 01 Rule lifecycle
- 02 MITRE mapping
- 03 Quality metrics

OPERATIONS

- 01 Case ownership
- 02 SLA tracking
- 03 Collaboration

PLATFORM

- 01 Automated tests
- 02 Observability
- 03 Resilience



The repository documents architecture, controls and deployment

FRONTEND


golden-horizon-soc-frontend

- React application
- Analyst workflows
- Dashboards + reports
- Vercel deployment

BACKEND

golden-horizon-soc-backend

- Express API
- PostgreSQL integration
- Socket.IO notifications
- Render deployment



Agent operations need a visible control plane


A product direction for governing agents, approvals, evaluations and operational risk.

OVERVIEW

The control tower makes agent behavior observable and governable


AgentOps extends security operations thinking to AI agents: every run has identity, policy, evidence and an accountable decision path.

- Inventory agents and owners
- Observe runs, tools and decisions
- Route high-risk actions for approval
- Evaluate quality and safety over time



ARCHITECTURE

A control plane surrounds every agent run



The stack is organized around control, evidence and integration

01

CONTROL PLANE

Policies • approvals • identity

02

RUNTIME

Agent orchestration • tool
adapters

03

EVIDENCE

Traces • logs • evaluations

04

INTEGRATION

Security data • tickets • cloud
APIs



FEATURES

Core features convert autonomy into accountable operations

GH
AGENTOPS
CONTROL CENTER

OVERVIEW

AGENTS

ACTIVITY

POLICIES

POLICY ALERTS

APPROVALS

AUDIT LOGS

SETTINGS

77% OPERATIONAL

RUNTIME EVENTS

RUNTIME ACTIVITY

Events generated by connected agents and runtime services

RUNTIME STREAM

EVENTS

ALL

HEARTBEATS

APPROVAL REQUESTS

VIOLATIONS

TIME	AGENT	EVENT	STATE	MODE	DETAILS
Jul 14, 08:33 PM	AI SOC Agent SOC-AGENT-001	POLICY VIOLATION DETECTED	open	Agent Runtime	modify policy / critical
Jul 14, 08:37 PM	AI SOC Agent SOC-AGENT-001	POLICY VIOLATION DETECTED	open	Agent Runtime	delete logs / high
Jul 14, 08:43 AM	DevOps Agent DEVOPS-AGENT-001	HEARTBEAT OPERATED	offline	Agent Runtime	Response time: ---
Jul 14, 08:47 AM	AI SOC Agent SOC-AGENT-001	HEARTBEAT OPERATED	healthy	Agent Runtime	Response time: 95 ms
Jul 14, 08:57 AM	AI SOC Agent SOC-AGENT-001	HEARTBEAT OPERATED	healthy	Agent Runtime	Response time: 95 ms
Jul 14, 09:22 AM	AI SOC Agent SOC-AGENT-001	HEARTBEAT OPERATED	healthy	Agent Runtime	Response time: 100 ms

AGENT INVENTORY Owner, purpose, model and approved tool scope

RUN OBSERVABILITY Trace inputs, reasoning artifacts, tools and outputs

APPROVAL ROUTING Risk-based human review before consequential action

EVALUATION Quality, safety and drift scorecards

CHALLENGES

Governance must stay proportional to risk

LOW

Read-only retrieval

Log + sample evaluation

MEDIUM

Draft or recommend

Policy checks + review

HIGH

External change

Explicit approval + rollback



Principle: governance grows with consequence and risk.

Every consequential action remains reviewable

QUALITY

Task correctness

SAFETY

Policy compliance

CONTROL

Approval coverage

OPERATIONS

Failure recovery

TIME	PERFORMED BY	AGENT	ACTION	PREVIOUS STATE	NEW STATE	DETAILS
Jul 15, 03:34 AM	Navid	AI SOC Agent SOC-AGENT-001	AGENT DISABLED	active	disabled	Recorded by AgentOps
Jul 15, 12:04 AM	Navid	HR Assistant HR-AGENT-001	AGENT DISABLED	active	disabled	Recorded by AgentOps
Jul 14, 09:35 PM	Navid	AI SOC Agent SOC-AGENT-001	VIOLATION RESOLVED	investigating	resolved	Blocked action attempt
Jul 14, 09:34 PM	Navid	AI SOC Agent SOC-AGENT-001	VIOLATION INVESTIGATING	open	investigating	Blocked action attempt
Jul 14, 09:23 PM	Navid	AI SOC Agent SOC-AGENT-001	VIOLATION RESOLVED	investigating	resolved	Blocked action attempt
Jul 14, 09:21 PM	Navid	AI SOC Agent SOC-AGENT-001	VIOLATION INVESTIGATING	open	investigating	Blocked action attempt

Target outcome

A reviewable operating record for every consequential agent action.

Approval workflows create a deliberate human boundary

1. INVENTORY

Register one agent, owner and approved scope.

2. OBSERVE


Capture runs, tool calls and outcomes.

3. GOVERN

Add policy gates and human approvals.

4. EVALUATE

Score quality, safety and recovery.



Repository strategy: ship the control plane as composable modules

/agents

Registry + ownership

/policies

Risk + approvals

/traces

Run evidence

/evaluations

Quality + safety

/integrations

SOC + cloud adapters

GH

AGENTOPS

CONTROL CENTER

OVERVIEW

AGENTS

ACTIVITY

POLICIES

POLICY ALERTS

APPROVALS

AGENT LOGS

SETTINGS

GOVERNANCE MONITOR

POLICY ALERTS

Governance violations detected across registered agents

TOTAL VIOLATIONS

2

OPEN

0

INVESTIGATING

0

RESOLVED

2

CRITICAL

1

VIOLATION GUIDES

ALL

OPEN

INVESTIGATING

RESOLVED

SEVERITY

AGENT

VIOLATION TYPE

ATTEMPTED ACTION

STATUS

DETECTED

ACTIONS

CRITICAL

AI SOC Agent
new agent 002

Blocked action attempt

modify_policy

RESOLVED

14.10.19.23 PM

VIEW DETAILS

HIGH

AI SOC Agent
new agent 002

Blocked action attempt

delete_logs

RESOLVED

14.10.19.21 PM

VIEW DETAILS

1

2

YOUR OPERATIONAL

GitHub profile: github.com/golden-horizon

The portfolio is backed by public repositories and working code

golden-horizon

Overview

Repositories 20

Projects

Packages

Stars 18

Type / to search

ghobadpour

on

I am an Australia-based DevOps and cloud security learner with a strong interest in secure cloud environments and automation.

agentic-ai-security-operations-platform

Public

AI-powered SOC Analyst Agent built with Python, Ollama, Qwen, MITRE ATT&CK, MCP, and modular agent architecture for automated security incident analysis.

Python 1

golden-horizon-soc-backend

Public

Backend API for a Security Operations Center platform built with Node.js, Express, PostgreSQL, JWT authentication, and real-time incident processing.

JavaScript 1

golden-horizon-soc-frontend

Public

Modern React-based SOC dashboard for security monitoring, incident investigation, threat visualization, and real-time analytics.

JavaScript 1

298 contributions in the last year

Contribution settings

2026

2025

2024

	Jul	Aug	Sep	Oct	Nov	Dec	Jan	Feb	Mar	Apr	May	Jun	Jul
Mon													
Wed													
Fri													

GitHub is the source of truth for the working portfolio



AI SOC REPOSITORY

github.com/golden-horizon/agent-ai-security-operations-platform



SOC FRONTEND

github.com/golden-horizon/golden-horizon-soc-frontend



GITHUB PROFILE

github.com/golden-horizon

Scan to explore the code and deployed experience.

CONTACT

Let's build security systems people can trust

NAVID GHOBADPOUR

AI • Cyber Security • Cloud Engineer

EMAIL

navid77g@gmail.com

LINKEDIN

[linkedin.com/in/navid-ghobadpour-90178180](https://www.linkedin.com/in/navid-ghobadpour-90178180)

GITHUB

github.com/golden-horizon

PHONE

Available on request



Sydney, NSW, Australia

END / 43